Shared memory and Consistency Models

Algorithmique répartie avancée - ARA
Master2

Luciana Arantes

Shared Memory

- **Shared Memory abstractions are programming abstractions that encapsulate read-writes forms of storage among processes**
  - *Motivation*: programming with shared memory model is considered easier than with message passing.
- **In shared-memory model, processes access concurrently data objects or memory location.**
- **Shared Memory variables are usually called** *read-write registers*

Shared memory mutual Exclusion

- **Mutual exclusion problem** [Taubenfeld06]
  - ➢ Two requirements should be satisfied:
    - ■ *Safety* : no two processes (threads) are in the critical section (CS) at the same time.
    - ■ *Deadlock- free*: if a process (thread) is trying to enter its CS, then some process (thread), eventually enters its CS.
      - ❑ Guarantees global progress property.
      - ❑ Does not prevent starvation.
  - ➢ Some algorithms provide a third property which ensures lack of starvation
    - ■ *Starvation-freedom*: if a process is trying to enter its CS, then this process must eventually enter its CS.

Remainder code
Enter CS code
Critical section CS
Exit CS code

Mutual exclusion: Peterson algorithm

- **Two processes**
- **The algorithm uses two shared variables: *flag[2]* and *turn*.**
  - **- A *flag* value of 1 indicates that the process wants to enter the CS.**
  - **- The variable *turn* holds the ID of the process whose turn it is.**
- **The algorithm is starvation-free**

Shared Variables:
boo flag[2] = {0,0}; int turn;

P_0:

```
flag[0] = 1; turn = 1;
while (flag[1] && turn==1);
// critical section
 ...
// end of critical section
flag[0] = 0 ;
```

P_1:

```
flag[1] = 1 ; turn = 0 ;
 while (flag[0] && turn == 0);
 // critical section
  ...
 // end of critical section
 flag[1] = 0;
```

Mutual exclusion: Lamport's Bakery Algorithm

- *N processes*
- **Starvation-free: satisfies mutual exclusion in first-come first-served**

Shared Variables:
  boo choosing[n];
  int timestamp[n];

Initialization:
  chosing[1..n] := 0;
  timestamp[1..n] := 0;

Entry CS Code:
  choosing [i] := 1;
  timestamp[i] := 1 + $\max_{1..n}$(timestamp[k]);
  chosing[i] := 0;
  for j := 1 to n do {
     await(chosing[j]=0);
     await (timestamp[j] == 0 **or** (timestam[i],i < timestamp[j],j)) ;

Exit Code:
timestamp[i] := 0;


Read-write Registers

- **A register is an abstraction of shared variable**
- **Implementation**
  - Provided by multiprocessors machine at hardware level
    - Array of hardware shared registers
  - Can also be implemented over processes that communicate through message passing and do not share any shared device [Guerraoui and Rodrigues 06], [Kshemkalayani and Singhal 08]
    - Shared memory emulation (distributed shared memory)

3

Read-write Registers

- **Store values that are accessed by *read* and *write* operations**
  - Process/Threads exchange information by invoking these operations
    - RW registers used for process/thread communication and synchronization

RW Registers [Lamport 86]

- **Definition 1: A RW register *x* is characterized by two operations:**
  - *write* $(x,v) \longrightarrow ok$: writes value *v* to register *x* and returns *ok*
  - **read**(x) $\longrightarrow v$: reads the register *x* and returns its value *v*.
- **Definition2 (Precedence) : for two operation o_1 and o_2, we say that:**
  - o_1 *precedes* o_2 whenever o_1 returns before o_2 is invoked (*sequential*)
  - o_1 is *concurrent* with o_2 when neither operation precedes the other one.

RW Registers [Lamport 86]

- **If a register is used by a single process, and we assume that there is no failure, we can define the following properties:**
  - **Safety**: Every read returns the *last* value written
  - **Liveness**: every operation eventually completes
- **Concurrency:**
  - In practice execution is not sequential
  - What is the meaning of "a read returns the last write" if both operations are concurrent?
    - It depends on the *semantics* of concurrent accesses offered by the register.
      - Safe, regular, and atomic registers.

Memory consistency model

- **Memory coherence is the ability of the system to execute memory operations correctly.**
  - Considering all the possible interleaving of operations issued by concurrent processes/threads, ensuring memory coherence becomes identifying which of these sequences of interleaving are correct.
  - Memory consistency model defines the sets of allowable memory access ordering.

Linearizibility

- **Consistency criteria for ordering concurrent accesses**
  - All operations appear to be executed atomically and sequentially
  - A global time scale needs to be simulated.
    - All processes (threads) need to agree on a common total order.

- **Other consistency model more relaxed**
  - Sequential consistency, causal consistency, PRAM, weak, etc
    - Discussed later

Linearizibility

- **For any concurrent execution, there is a total order of the operations such that each read to a location (variable) returns the value written by the last write to this location (variable) that precedes it in the order.**
- **This total order must be consistent with the temporal order of operations**
  - If one operation finishes before another begins, the former must precedes the latter in the total order.
    - Respect of the order of non overlapping operations.
  - For operations that overlap, all the processes/threads see the same ordering of events, which is equivalent to the global time occurrence of non overlapping events.
    - All operations appear to be executed atomically and sequentially.

Linearizibility

- **Each operation *op* (*Read* or *Write*) has an *invocation* and *response* events.**
  - ➢ An execution in global time is viewed as sequence *Seq* of such invocations and responses.
    - A *Seq* is linearizable is there is a permutation *Seq'* such that:
      - ❑ For every variable *v*, Seq'_v is such that each *Read* returns the most recent *Write* that immediately preceded it.
      - ❑ If the response of op_1 occurred before the invocation of op_2 in *Seq*, then op_1 occurs before op_2 en *Seq'*.

Linearizibility

Conflit avec la lecture de x, p2 retourne 0 alors que ça devrait être 1

Not linearizable!!

Linearizibility

P₁   write(x,1)   ok   read(x)   return (1)

P₂   read(x)   return (0)

overlapping

Sequential execution
read(x)   return (0)   write(x,1)   ok   read(x)   return (1)

Linearizable

Linearizibility

P₁   write(x,1)   ok   read(y)   return (2)

P₂   write(y,2)   ok   read(x)   return (1)

overlapping

Sequential execution
write(y,2)   ok   write(x,1)   ok   read(x)   return (1)   read(y)   return (2)
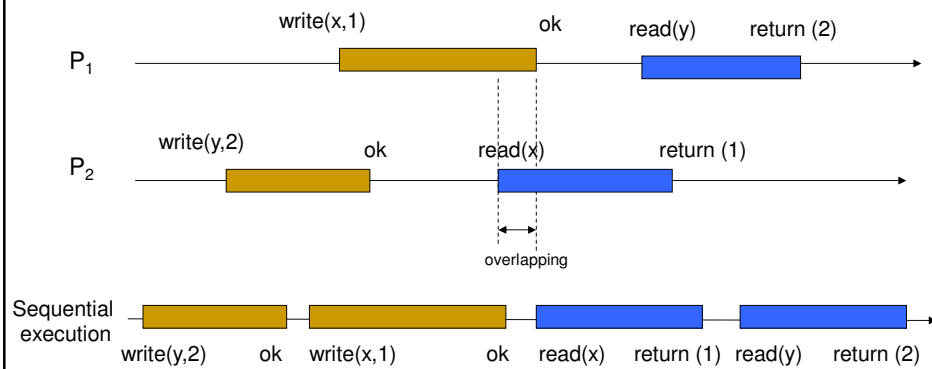
Linearizable

## Implementing Linearizibility on Message Passing

```
int x,

upon operation(op,val)   from application
   /* Read or Write */
   total_order_broadcast  (op, id_i);

upon reception of message <read, val, id>
   if (id = id_i)
      /* own request which was broadcast */
      return x;
```

```
upon reception of message  <write, val, id>
   x=val;
   if (id = id_i)
      /* own request which was broadcast */
      return ack to application
```

Reads must also participate in the total order broadcast

Not total order broadcast

Not linearlizable

Types of Registers

- **Semantics of *Read* and *Write* operations under concurrent accesses.**
  - In the face of concurrent *read* and *write* operations, the value returned by a *read* is unpredictable.
    - The order of access depends on the properties of the register
    - Implicit assumption of a global time
  - Three types of registers [Lamport 86]:
   - Safe
    - Regular
    - Atomic

Safe register

- *Read* **does not overlap with a** *write*
  - ➤ *Read* returns the most recently written value.
- *Read* **overlaps with a** *write*
 - ➤ *Read* returns any value that the register could possibly have.

Regular Register

- **It is a** *safe* **register and**
  - ➤ If *read* overlaps with a write
   - *Read* returns either the most recently value or a concurrently written value.

Atomic Register

- **It is a *regular* and**
  - *read* and *write* that overlap are *linearizable*
    - There exists an equivalent totally ordered sequential execution of them.

Characteristics of Registers

- **Semantics**

  24 types of registers

  - Safe, regular, atomic
- **Value**
  - Binary, integer
- **Write accesses**
  - Single-writer (SW), multi-writer (MW)
- **Read accesses**
 - Single-reader (SR), multi-reader (MR)

Register Construction

- **Design a more complex register using simpler registers.**

R   *write*

write	*write*		*write*
R_1	R_2	• • •	R_m
read	*read*		*read*

read

m individual registers

Construction of MRSW safe (regular) register with SRSW safe (regular) registers

- ***n* SRSW safe (regular) registers: $R_1, \ldots R_n$**
- **The single writer is process P_0 and the *n* readers are $P_1 \ldots P_n$.**
  - Multiple readers are not allowed to access the same SRSW safe (regular) register
    - A reader P_i can read only SRSW register R_i (the only reader)
      - Data must be replicated
  - P_0 can write to the *n* registers (the only writer)
    - P_0 writes the same value to the *n* registers

Construction of MRSW safe (regular) register with SRSW safe (regular) registers (cont.)

When a *read* by Pi and a *write* by P_0 do not overlap at Ri, the *read* returns the correct value; otherwise:

- *safe*: the *read* returns a legitimate value.

- *regular*: the *read* returns either the earlier value or the value being written.

SRSW safe (regular) registers $R_1 \dots R_n$

Single writer: P_0
multiple reader: $P_1, \dots P_n$

Write(val)
    for i=1 to n
        R_i = val.

Read (val)
  val = R_i
    return val


Construction of MRSW atomic register with SRSW atomic registers

- ***n* SRSW atomic registers: $R_1, \dots R_n$**
- **The single writer is process P_0 and the *n* readers are $P_1 \dots P_n$.**
- **The previous solution does not always ensure linearizable executions**

Non linearizable execution :
    • P_1 reads R_1 before the write
    • P_2 reads R_2 after the write

Construction of MRSW atomic register with SRSW atomic registers (cont.)

- **Solution [Israeli and Li 93]**
  - ➤ ***Read***: P_i must chose a value among R_i and the values that the other processes have last read;
    - *LastRead_val [n,n] of <data,seq>*: nxn SRSW atomic register that provides such information.
      - ❑ *LastRead_val [i,j]:* value of P_i's last returned read which was informed to P_j.
    - Before returning the value, the reader informs the other processes of the returned value.

Construction of MRSW atomic register with SRSW atomic registers

LastRead_val[n,n]

Construction of MRSW atomic register with SRSW atomic registers (cont.)

SRSW atomic registers of type <data,seq>: R_1 … R_n
SRSW atomic registers of type <data,seq>: LastRead_val[n,n]

Local Variables:
  int seq, j, latest;
<data,seq> LR[n+1]; /* last returned read value of other processes*/

Write(val)
  seq++;
  for j=1 **to** n
    R_i = <val,seq>.

Read (val)
  LR[0]=R_i
  for j = 1 **to** n
    /* get latest value stored for P_i by P_j */
    LR[j]=LastRead_val [j,i]

  find max **such that for all** latest<> k
    LR[latest].seq >= LR[k].seq;

  for j = 1 **to** n
    LastRead_val [i,j] = LR[latest];

  val = LR[latest]
  return (val)

Construction of MRSW regular register on message passing system with crashes

- **Emulation of a MRSW regular register**
  - One specific process P_0 can invoke a *write* operation and any other can invoke a *read* operation on the register.
- **Use of a perfect failure detector**
- **Each process stores a copy of the current register value in a variable that is local to a process**
- **Read-one Write-all algorithm:**
  - The writer updates the value of all processes which it does not detect faulty.
    - All processes acknowledge the receipt of the new value.
    - *Write* completes when all acknowledges from correct process is received.
  - The reader just return the value stored locally.

Construction of MRSW regular on message passing system with crashes

Local Variables:
writeSet= Ø;
reg = 0;
correct = Π

Read ( )
return reg;

Write(val)
Bebbroadcast <val>

Upon event crash <j>
correct = correct / {j}

Upon event Bebdelivery <val,j>
reg = val;
send <j,ack>

Upon event reception <ack,j>
writeSet = writeSet U {j}

Upon exist r **such that** correct C writeSet
writeSet= Ø;
return <ok>;

Atomic Operations

- **Read, write**
- **Examples of other atomic operations:**
  - *test-and-set (r:register; val:value): value*
    - The value *val* is assigned to *r,* and the old value of *r* is returned.
  - *read-modify-write (r:register; f:function):value*
    - The value of *f(r)* is assigned to *r*, and the old value of *r* is returned.
  - *compare-and-swap (r:register; key, new :value):value*
    - If the current value of *r* is equal to *key*, then the value of *r* is set to *new* ; otherwise *r* is not changed. The old value of *r* is returned.
  - *fetch-and-add (r:register; val:value): value*
    - The value of *r* is incremented by *val* and the old value of *r* is returned
  - *enqueue (Q:queue,val:value) ; dequeue (Q:queue):value*
    - Operations for a FIFO queue object.
  - ...

Mutual exclusion using test-and-set

• *N processes*
• **Not starvation-free**

```
function T&S (r:register, val:value): value;
   /* atomic*/
   temp = r;
   r=val;
   return (temp);
```

Shared Variables:
  register reg =false;

Local Variable:
  blocked;

Entry CS Code:
  blocked=true;

  repeat
      blocked=T&S(reg,blocked);
  until blocked=false;   *false*

Exit Code:
  reg=false;

Shared Atomic Objects

- **A shared atomic abject is a data structure exporting a set of operations that can be invoked concurrently by the processes (threads) of the system.**
  - Each object has a type which defines the set of operations (primitives, methods) that the object supports.
    - Object is accessed only by using such operations
    - Each object has sequential specification that specifies how the object behaves when these operations are applied atomically.
  - There are objects which have more synchronization power than atomic *Read/Write* registers.

Examples of atomic shared objects

- **Registers**
- **Test-and-Set object**
  - A shared register that supports *write* and *test-and-set* operations.
- **Read-modify-write object**
  - A shared register that supports *read-modify-write* operation.
- **Compare-and-swap object**
  - A shared register that supports *compare-and-swap* operation.
- **Queue**
  - A shared register that supports *enqueue* and *dequeue* operations.
- **….**

Registers : Failure Issues

- **If a process (thread) can fail by crashing, the operation invoked by it might not complete.**
  - If a process (thread) invokes a *write* and crashes, the *write* is considered to be concurrent with any *read* that did not precede it.
- **Any process (thread) that invokes a *read* or *write* operation and does not crash eventually returns from this invocation.**

Synchronization of operations

- **Blocking operations**
  - **A delay or crash of a process (thread) can prevent others from making progress.**
    - **e.g. Mutex, producer-consumer, etc.**
- **Non blocking operations**
  - **A delay or crash of a process (thread) can not prevent others from making progress.**
    - **Processes (threads) competing for a shared resource do not have their execution indefinitely postponed by other processes.**

Non blocking operations

- **Wait-freedom**
  - An operation on a shared object is *wait-free* if every invocation of the operation completes in a finite number of steps regardless of the number of steps taken by any other process.
    - A concurrent object is *wait-free* if all its operation is *wait-free*
    - Wait-freedom provides robustness and ensures per-process (per-thread) progress.
    - A process (thread) does not depend on other process (thread) , and its execution is does *wait-free*.
    - If n is the number of processes (thread) , $n-1$ processes (threads) can crash
      - A wait-free algorithm in a system with n processes (threads) is a *(n-1)-crash resilient* algorithm.

Wait-free shared memory consensus

- **Consensus is impossible in an asynchronous shared memory system in the crash failure model.**
  - Extended from the impossibility in message-passing (FLP)
    - Shared memory can be emulated by message passing
  - In the face of a potential crash it is not possible to distinguish between a crashed process and a process which extremely slow in doing its *Read* or *Write* operation.

Consensus impossible dans un système avec passage de messages asynchrone

Wait-free shared memory consensus

- **There are objects for which there is a *wait-free* algorithm for reaching consensus in a *n-processes* system.**
  - Objects that provide stronger synchronization than safe, regular and atomic objects.
    - e.g. test-and-set objects, compare-and-swap objects, FIFO queue objects, etc.

Consensus Number

- An object of type *X* has consensus number *k*, if *k* is the largest number for which the object *X* can solve wait-free *k*-process consensus in an asynchronous system subject to *k-1* crash failures, using only objects of type *X* and *read/write* objects.

Consensus number of some objects

(∞)	**Compare&Swap**
	• • •
(3)	
(2)	**FIFO Queue** **Test&Set** **Fetch&Add**
(1)	**Register**

Wait-free consensus : Test-and-set object

• Two processes

x_i: initial choice

Shared Variables :
test-and-set reg = $\perp$;
r/w reg choice[2] ={$\perp$, $\perp$};

Local Variables:
int val;

choice[i] =x_i;
val= T&S(reg,x_i)
if (val = $\perp$)
decide (x_i)
else
decide (choice[i-1]);

choice[0] = x_0
T&S(reg,x_0)
P$_0$ — decide (x_0)

choice[1] = x_1
T&S(reg,x_1)
P$_1$ — decide (x_0)

It does not work with more than 2 processes:
which value to chose from *choice[ ]* ?

Wait-free consensus : FIFO Queue object

• Two processes

x_i: initial choice

Shared Variables :
queue reg Q= <0> /* Q initializé */
r/w reg choice[2] ={$\perp$, $\perp$};

Local Variables:
int temp;

choice[i] =x_i;
temp= dequeue(Q) ;
if (temp == $\perp$)
/* queue empty*/
decide (choice[i-1]);
else
decide (x_i)

choice[0] = x_0
dequeue(Q)
P$_0$ — decide (x_0)

choice[1] = x_1
dequeue(Q)
P$_1$ — decide (x_0)

Wait-free consensus : Compare-and-set object

- ∞ **processes**

x_i: **initial choice**

Shared Variables :
compare-and-swap reg;

Local Variables:
**int** temp;

temp= C&S(reg, $\perp$, x_i);
if (temp = $\perp$)
  decide (x_i)
else
    decide (temp);

function C&S (r:register; key,new: value): value;
  /* atomic*/
  temp = reg;
  if (key=temp) reg=new;
  return (temp);

P_0 C&S(r, $\perp$,x_0)   decide (x_0)

P_1 C&S(r, $\perp$,x_1)   decide (x_0)

P_2 C&S(r, $\perp$,x_2)   decide (x_0)

P_3 • • •

Non-blocking operations

- **Wait-freedom**
  - It is the strongest non-blocking guarantee of progress, combining guaranteed system-wide throughput with starvation-freedom. An algorithm is *wait-free* if every operation has a bound on the number of steps the algorithm will take before the operation completes.
- **Lock-freedom**
  - Lock-freedom allows individual threads to starve but guarantees system-wide throughput. An algorithm is lock-free if every step taken achieves global progress (for some sensible definition of progress). All wait-free algorithms are lock-free.
- **Obstruction-freedom**
  - An algorithm is obstruction-free if at any point, a single thread executed in isolation (i.e., with all obstructing threads suspended) for a bounded number of steps will complete its operation. All lock-free algorithms are obstruction-free.

Memory consistency Models

- **Linearizability = Strict or atomic consistency:**
- **Other models more relaxed :**
  - Sequential consistency
  - Causal consistency
  - PRAM consistency
  - Consistency models based on synchronization variables
    - Weak
    - Entry consistency
    - Release
    - Lazy release

Sequential Consistency

- The result of any execution is the same as if all operations of the processors were executed in some sequential order.
- The operations of each individual processor appear in this sequence in the local program order.

• Any interleaving of the operations from the different processes is possible. But all processors must see the same interleaving.
• Even if two operations from different processes (on the same or different variables) do not overlap in a global time scale, they may appear in reverse order in the common sequential order seen by all.

Sequential Consistency

write(x,1)   ok   read(x)   return (1)

P₀ ———————[====]—————————————————————[====]—————————→

                        read(x)   return (0)

P₁ ———————————————————————[====]—————————————————————→

Sequential
execution ———[====]————————[====]————————[====]——————→
           read(x) return (0) write(x,1) ok   read(x) return (1)

Sequentially consistency !

08/11/2012   ARA: Shared Memory   49

Sequential Consistency

write(x,1)   ok   read(x)   return (0)

P₀ ———————[====]—————————————————————[====]—————————→

                        read(x)   return (0)

P₁ ———————————————————————[====]—————————————————————→

Not Sequential consistency !

08/11/2012   ARA: Shared Memory   50

25

Causal Consistency

- **Only *writes* that are *causally related* must be seen by all processes in the same order. Concurrent writes must be seen in a different order.**
 - The causal relation is defined as:
    - Local order of events of a processes define local causal order.
    - A *write* operation causally precedes a *read* operation of another process if the *read* returns the value written by the *write* operation.
    - The transitive operation of the above two relations defines the global causal order.

Causal Consistency

Not causal consistency !

Causal Consistency

P0 write(x,1) ok

P1 write(x,2) ok

P3 read(x) return (1) read(x) return (2)

P4 read(x) return (2) read(x) return (1)

Causal consistency !
The writes are concurrent

Causal Consistency

P0 write(x,1) ok

P1 read(x) return (1) write(y,2) ok

P2 read(y) return (2) read(x) return (0)

Not causal consistency !

PRAM Consistency

- *Writes* **done by a single processor are received by all other processes in the order in which they were issued but** *writes* **from different processes may be seen in a different order by different processes**
  - Only the local causality relation needs to be seen by other processes.
    - Writes from the same processes mustbe seen by the others in order they were issued.

08/11/2012   ARA: Shared Memory   55


PRAM Consistency

PRAM consistency !

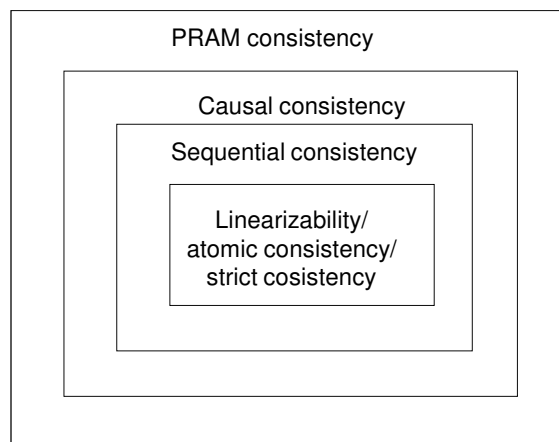08/11/2012   ARA: Shared Memory   56

28

PRAM Consistency

write(x,1)   ok   write(y,2)   ok

P_0 ――[   ]――――[   ]――――――――――――――――――→

                              read(y)   return (2)   read(x)   return (0)

P_1 ――――――――――――――[   ]―――[   ]――→

Not PRAM consistency !

Hierarchy of consistency models

PRAM consistency

Causal consistency

Sequential consistency

Linearizability/
atomic consistency/
strict cosistency

29

Consistency models: applications

- **Memory consistency models can be applied to other domains**
  - Example: data base systems where data are replicated for fault tolerance and performance reasons on several servers.
    - Clients: c_1, c_2,...
    - Servers: s_1, ..., s_n
    - Operations at client side: **read** and **write**
      - *read(x)* returns the value of data item x
      - *write(x,val)* updates the value of x with *val* and returns an acknowledgement
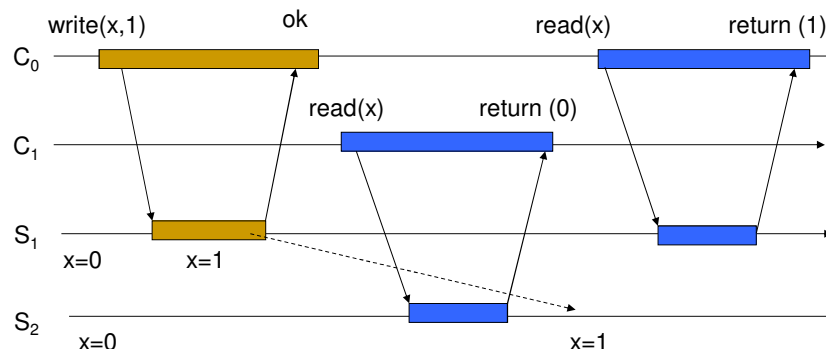   - An update from a client is broadcast to all servers

Multiple clients and multiple servers

Is it inconsistent?

Multiple clients and multiple servers
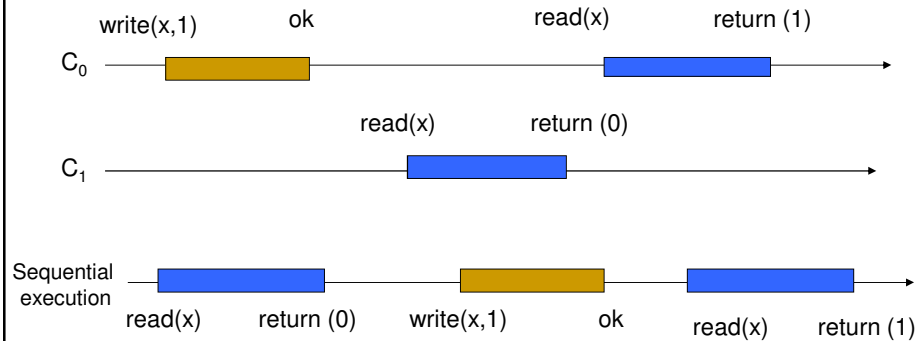
Yes if sequential consistency provided

Software Transactional Memory (STM)

- **A transaction is a finite sequence of local and shared memory instructions:**
  - *Read_transcational*: reads the value of a shared location into a local register
  - *Write_transactional*: stores the contents of a local register into a shared location.
- **The data set of a transaction is the set of shared locations accessed by these instructions.**
- **Any transaction may either fail or complete successfully.**
  - Changes are visible atomically to other processes.
- **A STM implementation is non-blocking**
  - It is *wait-free* if any process which repeatedly executes the transaction terminates successfully after a finite number of attempts.

Software Transactional Memory (STM)

- **A software transactional memory (STM) is a shared object which behaves like a memory that supports multiple changes to its addresses by means of transactions.**
  - It is a concurrency control mechanism analogous to database transactions for controlling access to shared memory.
  - Any implementation of it should satisfies the following property (atomicity and linearizibility (serializability)):
    - Each transaction appears to execute instantaneously and after transactions are executed concurrently, the state of the shared data and the outcome of the transactions are the same as if these transactions were executed serially in some order.

Software Transactional Memory (STM)

- **Example of transaction**
  - A k-word Compare&Swap Transaction

```
k_word_C&S (N, DataSet [ ], key [ ], New [ ]) {
BeginTransaction

  for i = 1 to k do
    if Read_transactional (DataSet[i] != key[i])
       return C&SFailure
  for i = 1 to k
       Write_transactional (DataSet[i], New[i])
  return C&SSuccess

EndTransaction
}
```

Bibliography

- G. Taubenfeld, *Synchronization Algorithms And Concurrent Programming*, Pearson Prentice Hall, 2006.
- M. Herlihy, and N. Shavit, *The Art of Multiprocessor Programming,* Morgan Kauman Publishers*, 2008.*
- A. D. Kshemkalyani, and M. Singhal, *Distributed Computing: principles, algorithms, and systems*, Cambridge University Press, 2008.
- R. Guerraoui, and L. Rodrigues. *Reliable Distributed Programming*, Springer, 2006.
- M. Herlihy. Wait-free synchronization. *ACM Transaction on Programming Languages and Systems*, 13(1), 1991, pages 124-149.
- Nancy Lynch, *Distributed Algorithms*, Morgan Kaufman Publishers, 1996.

Bibliography (cont.)

- L. Lamport. On interprocess communication, *Distributed computing*, 1(2),1986, pages 77-85.
- L. Lamport, A new solution of Dijkstra's concurrent programming problem, *Communication of the ACM*, 17(8), 1974, pages 453-455.
- A. Israeli, and M. Li. Bounded timestamps. *Distributed Computing*, 6(4), 1993, pages 205-209.
- M. Ahamad, G. Neiger, J. E. Burns, P. Kohli, and P. W. Hutto Causal memory: definitions, implementation, and programming. *Distributed Computing*, 9(1), 1995, pages 37-49
- R. Lipton and J. Sandberg. *PRAM: a Scalable Shared Memory*. Technical Report CS-TR-180-88, Princeton University, 1988.
- N. Shavit, and D. Touitou. Software Transactional Memory. *PODC* 1995, pages 204-213.